Coupled Markov chains with applications to Approximate Bayesian Computation for model based clustering

E. Bertoni, M. Caldarini, F. Di Filippo, G. Gabrielli, E. Musiari

11 november 2021

A complex problem

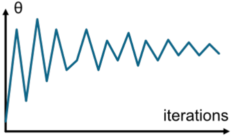

A complex problem

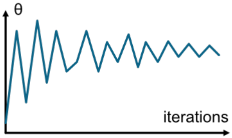

likelihood

A complex problem

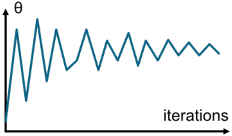

A complex problem

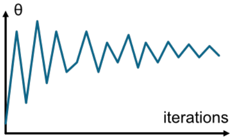

likelihood intractable

⇓

Unbiased Markov chain Monte Carlo methods with couplings

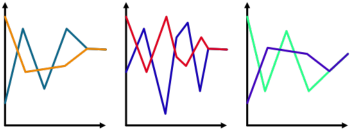

A complex problem

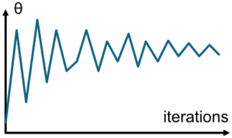

⇓

Unbiased Markov chain Monte Carlo methods with couplings

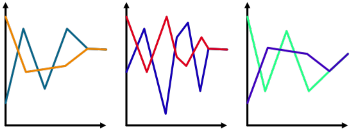

⇓

Approximate Bayesian Computation

Unbiased Markov chain Monte Carlo methods with couplings

E. Bertoni, M. Caldarini, F. Di Filippo, G. Gabrielli, E. Musiari

The road to parallelization: coupling of Markov chains

Faster MCMC $\Longrightarrow$ Parallelization

The road to parallelization: coupling of Markov chains

Faster MCMC $\Longrightarrow$ Parallelization $\Longleftrightarrow$ **Unbiased estimator**

The road to parallelization: coupling of Markov chains

Faster MCMC $\Longrightarrow$ Parallelization $\Longleftrightarrow$ **Unbiased estimator**

Exact estimations algorithms using **coupling of Markov chain**.

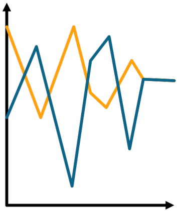

Rhee–Glynn estimator I

The goal is to estimate

$$\mathbb{E}_\pi[h(X)] = \int h(x)\pi(\mathrm{d}x).$$

The estimator we are going to construct is based on a coupled pair of Markov chains, $(X_t)_{t\geq 0}$ and $(Y_t)_{t\geq}$, which marginally start from π_0 and evolve accordingly to P.

Rhee–Glynn estimator II

We consider some assumptions:

1. as $t \to \infty$,
$$\mathbb{E}[h(X_t)] \to \mathbb{E}_\pi[h(X)];$$
and there exists $\eta > 0$ and $D < \infty$ such that $\mathbb{E}[|h(X_t)|^{2+\eta}] \leq D$ for all $t \geq 0$;

Rhee–Glynn estimator II

We consider some assumptions:

1. as $t \to \infty$,
$$\mathbb{E}[h(X_t)] \to \mathbb{E}_\pi[h(X)];$$
and there exists $\eta > 0$ and $D < \infty$ such that $\mathbb{E}[|h(X_t)|^{2+\eta}] \leq D$ for all $t \geq 0$;
2. the chains are such that the meeting time
$$\tau = \inf\{t \geq 1 : X_t = Y_{t-1}\}$$
satisfies $\mathbb{P}(\tau > t) \leq C\delta^t$ for all $t \geq 0$, for some constants $C < \infty$ and $\delta \in (0, 1)$;

Rhee–Glynn estimator II

We consider some assumptions:

1. as $t \to \infty$,
$$\mathbb{E}[h(X_t)] \to \mathbb{E}_\pi[h(X)];$$
and there exists $\eta > 0$ and $D < \infty$ such that $\mathbb{E}[|h(X_t)|^{2+\eta}] \leq D$ for all $t \geq 0$;
2. the chains are such that the meeting time
$$\tau = \inf\{t \geq 1 : X_t = Y_{t-1}\}$$
satisfies $\mathbb{P}(\tau > t) \leq C\delta^t$ for all $t \geq 0$, for some constants $C < \infty$ and $\delta \in (0, 1)$;
3. the chains stay together after meeting:
$$X_t = Y_{t-1} \text{ for all } t \geq \tau.$$

Rhee–Glynn estimator III

Thanks to the previous assumptions we can prove that:

$$\mathbb{E}_{\pi}[h(X)] = \mathbb{E}[h(X_k) + \sum_{t=k+1}^{\tau-1} \{h(X_t) - h(Y_{t-1})\}];$$

Rhee–Glynn estimator III

Thanks to the previous assumptions we can prove that:

$$\mathbb{E}_\pi[h(X)] = \mathbb{E}\Big[h(X_k) + \sum_{t=k+1}^{\tau-1} \{h(X_t) - h(Y_{t-1})\}\Big];$$

and we define the Rhee–Glynn estimator as:

$$H_k(X, Y) = h(X_k) + \sum_{t=k+1}^{\tau-1} \{h(X_t) - h(Y_{t-1})\}$$

which is **unbiased** by construction.

Time-averaged estimator I

$$H_{k:m}(X,Y) = \frac{1}{m-k+1}\sum_{l=k}^{m} h(X_l) + \sum_{l=k+1}^{\tau-1} \min(1, \frac{l-k}{m-k+1})\{h(X_l) - h(Y_{l-1})\}$$

Time-averaged estimator I

$$H_{k:m}(X, Y) = \underbrace{\frac{1}{m-k+1}\sum_{l=k}^{m} h(X_l)}_{MCMC_{k:m}} + \sum_{l=k+1}^{\tau-1} \min(1, \frac{l-k}{m-k+1})\{h(X_l) - h(Y_{l-1})\}$$

- $MCMC_{k:m}$ is the standard MCMC average;

Time-averaged estimator I

$$H_{k:m}(X,Y) = \underbrace{\frac{1}{m-k+1}\sum_{l=k}^{m} h(X_l)}_{MCMC_{k:m}} + \underbrace{\sum_{l=k+1}^{\tau-1} \min\left(1, \frac{l-k}{m-k+1}\right)\{h(X_l) - h(Y_{l-1})\}}_{BC_{k:m}}$$

- $MCMC_{k:m}$ is the standard MCMC average;
- $BC_{k:m}$ is the bias correction;

Time-averaged estimator II

1. draw X_0 and Y_0 from an initial distribution π_0 and draw $X_1 \sim P(X_0, \cdot)$;
2. set $t = 1$: while $t < \max\{m, \tau\}$ and:
   a. draw $(X_{t+1}, Y_t) \sim \bar{P}\{(X_t, Y_{t-1}), \cdot\}$;
   b. set $t \leftarrow t + 1$;
3. compute

$$H_{k:m}(X, Y)$$

with the time-averaged estimator.

Time-averaged estimator III

Metropolis–Hasting algorithm allow us to calculate the coupled kernel $\bar{P}\{(X_t, Y_{t-1}), \cdot\}$:

1. sample $(X^\star, Y^\star)|(X_t, Y_{t-1})$ from a maximal coupling of $q(X_t, \cdot)$ and $q(Y_{t-1}, \cdot)$;
2. sample $U \sim \mathcal{U}([0,1])$;
3. if
$$U \leq \min\left\{1, \frac{\pi(X^\star)q(X^\star, X_t)}{\pi(X_t)q(X_t, X^\star)}\right\}$$
then $X_{t+1} = X^\star$; otherwise $X_t = X_{t-1}$;
4. if
$$U \leq \min\left\{1, \frac{\pi(Y^\star)q(Y^\star, Y_t)}{\pi(Y_t)q(Y_t, Y^\star)}\right\}$$
then $Y_{t+1} = Y^\star$; otherwise $Y_t = Y_{t-1}$.

Time-averaged estimator III

Metropolis–Hasting algorithm allow us to calculate the coupled kernel $\bar{P}\{(X_t, Y_{t-1}), \cdot\}$:

1. sample $(X^\star, Y^\star)|(X_t, Y_{t-1})$ from a maximal coupling of $q(X_t, \cdot)$ and $q(Y_{t-1}, \cdot)$;
2. sample $U \sim \mathcal{U}([0, 1])$;
3. if
$$U \leq \min\left\{1, \frac{\pi(X^\star)q(X^\star, X_t)}{\pi(X_t)q(X_t, X^\star)}\right\}$$
then $X_{t+1} = X^\star$; otherwise $X_t = X_{t-1}$;
4. if
$$U \leq \min\left\{1, \frac{\pi(Y^\star)q(Y^\star, Y_t)}{\pi(Y_t)q(Y_t, Y^\star)}\right\}$$
then $Y_{t+1} = Y^\star$; otherwise $Y_t = Y_{t-1}$.

Approximate Bayesian Computation

E. Bertoni, M. Caldarini, F. Di Filippo, G. Gabrielli, E. Musiari

POLITECNICO MILANO 1863

Likelihood-free rejection sampling algorithm I

Inputs:

- a target posterior density $\pi(\theta|y_{obs}) \propto p(y_{obs}|\theta)\pi(\theta)$, consisting of a prior distribution $\pi(\theta)$ and a procedure of generating data under the model $p(y_{obs}|\theta)$;
- a proposal density $g(\theta)$, with $g(\theta) > 0$ if $\pi(\theta|y_{obs}) > 0$;
- an integer $N > 0$.

Likelihood-free rejection sampling algorithm I

Inputs:

- a target posterior density $\pi(\theta|y_{obs}) \propto p(y_{obs}|\theta)\pi(\theta)$, consisting of a prior distribution $\pi(\theta)$ and a procedure of generating data under the model $p(y_{obs}|\theta)$;
- a proposal density $g(\theta)$, with $g(\theta) > 0$ if $\pi(\theta|y_{obs}) > 0$;
- an integer $N > 0$.

Sampling for $i = 1, ..., N$:

1. generate $\theta^{(i)} \sim g(\theta)$ from sampling density g;
2. generate $y \sim p(y|\theta^{(i)})$ from the likelihood;
3. if $y = y_{obs}$, then accept $\theta^{(i)}$ with probability $\frac{\pi(\theta^{(i)})}{Kg(\theta^{(i)})}$, where $K \geq \max_\theta \frac{\pi(\theta)}{g(\theta)}$; else go to 1.

Likelihood-free rejection sampling algorithm I

Inputs:

- a target posterior density $\pi(\theta|y_{obs}) \propto p(y_{obs}|\theta)\pi(\theta)$, consisting of a prior distribution $\pi(\theta)$ and a procedure of generating data under the model $p(y_{obs}|\theta)$;
- a proposal density $g(\theta)$, with $g(\theta) > 0$ if $\pi(\theta|y_{obs}) > 0$;
- an integer $N > 0$.

Sampling for $i = 1, ..., N$:

1. generate $\theta^{(i)} \sim g(\theta)$ from sampling density g;
2. generate $y \sim p(y|\theta^{(i)})$ from the likelihood;
3. if $y = y_{obs}$, then accept $\theta^{(i)}$ with probability $\frac{\pi(\theta^{(i)})}{Kg(\theta^{(i)})}$, where $K \geq \max_\theta \frac{\pi(\theta)}{g(\theta)}$; else go to 1.

Output:

- a set of parameter vectors $\theta^{(1)}, ..., \theta^{(N)}$ which are samples from $\pi(\theta|y_{obs})$.

Likelihood-free rejection sampling algorithm II

Is this an efficient method for complex analysis?

Likelihood-free rejection sampling algorithm II

Is this an efficient method for complex analysis?

3 If $\| y - y_{obs} \| \leq h$, then accept $\theta^{(i)}$ with probability $\frac{\pi(\theta^{(i)})}{Kg(\theta^{(i)})}$, where $K \geq \max_\theta \frac{\pi(\theta)}{g(\theta)}$; else go to 1.

Approximate Bayesian Computation

$$\pi(\theta, y|y_{obs}) \propto \mathbb{1}(\| y - y_{obs} \| \leq h) p(y|\theta)\pi(\theta)$$

$$\Downarrow$$

$$\pi_{ABC}(\theta, y|y_{obs}) \propto K_h(u) p(y|\theta)\pi(\theta)$$

Approximate Bayesian Computation

$$\pi(\theta, y|y_{obs}) \propto \mathbb{I}(\| y - y_{obs} \| \leq h) p(y|\theta)\pi(\theta)$$

$$\Downarrow$$

$$\pi_{ABC}(\theta, y|y_{obs}) \propto K_h(u) p(y|\theta)\pi(\theta)$$

Where we used a **standard smoothing kernel function**:

$$K_h(u) = \frac{1}{h} K\left(\frac{u}{h}\right), \quad \text{with } u = \| y - y_{obs} \|$$

Summary statistics I

Is this feasible in practice?

Summary statistics I

Is this feasible in practice?

No, it's difficult to have $y \approx y_{obs}$: we should use a large h, obtaining a poor posterior approximation!

$\Longrightarrow$ use summary statistics $s = S(y)$

Summary statistics II

Critical decision: choice of summary statistics

Summary statistics II

Critical decision: choice of summary statistics Dimension of summary

statistics:

- large enough to contain as much as information about observed data as possible
- low enough to avoid curse of dimensionality of matching s and s_{obs}

$\Longrightarrow$ choose sufficient statistics, such that:

$$\pi(\theta|s_{obs}) \equiv \pi(\theta|y_{obs})$$

Distance measure

Distance measure: substantial impact on ABC algorithm efficiency

$$\| s - s_{obs} \| = (s - s_{obs})^\top \Sigma^{-1} (s - s_{obs})$$

Distance measure

Distance measure: substantial impact on ABC algorithm efficiency

$$\| s - s_{obs} \| = (s - s_{obs})^\top \Sigma^{-1} (s - s_{obs})$$

- $\Sigma =$ identity matrix $\rightarrow$ Euclidean distance
- $\Sigma =$ diagonal matrix of non-zero weights $\rightarrow$ Weighted Euclidean distance
- $\Sigma =$ full covariance matrix of s $\rightarrow$ Mahalanobis distance

ABC rejection sampling algorithm

Inputs:

- a target posterior density $\pi(\theta|y_{obs}) \propto p(y_{obs}|\theta)\pi(\theta)$, consisting of a prior distribution $\pi(\theta)$ and a procedure of generating data under the model $p(y_{obs}|\theta)$;
- a proposal density $g(\theta)$, with $g(\theta) > 0$ if $\pi(\theta|y_{obs}) > 0$;
- an integer $N > 0$;
- a kernel function $K_h(u)$ and a scale parameter $h > 0$;
- a low dimensional vector of summary statistics $s = S(y)$.

ABC rejection sampling algorithm

Sampling for $i = 1, ..., N$:

1. generate $\theta^{(i)} \sim g(\theta)$ from sampling density g;
2. generate $y \sim p(y|\theta^{(i)})$ from the likelihood;
3. compute summary statistic $s = S(y)$;
4. accept $\theta^{(i)}$ with probability $\frac{K_h(\|s - s_{obs}\|)\pi(\theta^{(i)})}{Kg(\theta^{(i)})}$, where $K \geq K_h(0) \max_\theta \frac{\pi(\theta)}{g(\theta)}$; else go to 1.

Output:

- a set of parameter vectors $\theta^{(1)}, ..., \theta^{(N)} \sim \pi_{ABC}(\theta|S_{obs})$.

Conclusions

Next steps

Our focus till now was to understand the fundamental concepts and collect the missing information.

The next step will be a **simple and separate implementation** of both solution to be tested on simulated data.

Further steps will consider the **integration** of both solution into a single implementation and the testing on more complex data.

Bibliography

Pierre Jacob, John O'Leary, and Yves Atchadé.
Unbiased markov chain monte carlo with couplings.
Journal of the Royal Statistical Society: Series B (Statistical Methodology), 82, 08 2017.

Peter W. Glynn and Chang han Rhee.
Exact estimation for markov chain equilibrium expectations, 2014.

Jeffrey S. Rosenthal.
Faithful couplings of markov chains: Now equals forever.
Advances in Applied Mathematics, 18(3):372–381, 1997.

Dylan Cordaro.
Markov chain and coupling from the past.
2017.

Jinming Zhang.
Markov chains, mixing times and coupling methods with an application in social learning.
2020.

S. A. Sisson, Y. Fan, and M. A. Beaumont.
Overview of approximate bayesian computation, 2018.

Y. Fan and S. A. Sisson.
Abc samplers, 2018.

Dennis Prangle.
Summary statistics in approximate bayesian computation, 2015.